

Assignment 1 - Image Deblurring

Name: Brian Schweigler

Matriculation number: 16-102-071

```
In [2]: import numpy as np
import scipy
import scipy.sparse
from scipy.sparse.linalg import inv
from PIL import Image
import matplotlib.pyplot as plt
from hessian_matrix import hessian_matrix
```

Test Image

You can use the following function to create a synthetic image:


```
In [3]: def create_random_binary_img_1():
    # randint returns an integer from 1 (inclusive) to 2 (exclusive), meaning it will al
    u = np.random.randint(1, 2, 30 * 30).reshape(30, 30).astype(float)
    u[5:9, 4:] = 0.5
    u[3:5, 4:] = 0
    u[10:15, :10] = 0
    u[10:15, 15:] = 0.8
    u[15:, 18:27] = 0.3

    for i in range(10):
        u[i + 17, i + 4:i + 8] = 0

    u[20:25:, 20:23] = 1
    u[20:28, 4] = 0
    u[20:28, 3] = 0.3
    u[20:28, 2] = 0.7

    return u

sample_image = create_random_binary_img_1()
plt.imshow(sample_image, cmap="gray")
plt.show()
```



hessian_matrix() function

We provide the implementation of the hessian matrix for $k = 0, 1, 2, 3$.

input variables:

- u: your image
- reg: regularization parameter
- k: one of the four kernel types 0-3

Example usage:

```
In [4]: u_ex = np.zeros(9).reshape(3, 3)
reg_ex = 1
k_type_ex = 3
A = hessian_matrix(u_ex, reg_ex, k_type_ex)
A_n_t = A.toarray()
print(A_n_t)
```

```
[[ 4.  -2.   0.  -2.   0.   0.   0.   0.   0. ]
 [-2.   6.5 -2.   0.5 -2.   0.   0.   0.   0. ]
 [ 0.  -2.   4.5  0.   0.5 -2.   0.   0.   0. ]
 [-2.   0.5  0.   6.5 -2.   0.  -2.   0.   0. ]
 [ 0.  -2.   0.5 -2.   9.  -2.   0.5 -2.   0. ]
 [ 0.   0.  -2.   0.  -2.   6.5  0.   0.5 -2. ]
 [ 0.   0.   0.  -2.   0.5  0.   4.5 -2.   0. ]
 [ 0.   0.   0.   0.  -2.   0.5 -2.   6.5 -2. ]
 [ 0.   0.   0.   0.   0.  -2.   0.  -2.   4. ]]
```

Implementation

```

In [5]: def derivative_img_term(g, u):
        """
        Derivative of the image term
        g: Blurred image
        u: The image
        """
        result = np.zeros(u.shape)
        for x in range(u.shape[0]):
            for y in range(u.shape[1]):
                # Corners
                if x == 0 and y == 0:
                    result[x, y] = 0
                elif x == 0 and y == u.shape[1] - 1:
                    result[x, y] = -g[0, y - 1] + 0.5 * u[1, y] + 0.5 * u[0, y]
                elif x == u.shape[0] - 1 and y == 0:
                    result[x, y] = 0
                elif x == u.shape[0] - 1 and y == u.shape[1] - 1:
                    result[x, y] = -g[x - 1, y - 1] + 0.5 * u[x, y] + 0.5 * u[x - 1, y]
                # Edges
                elif x == 0: # Upper edge
                    result[x, y] = -g[0, y - 1] + 0.5 * u[0, y] + 0.5 * u[1, y]
                elif x == u.shape[0] - 1: # Lower edge
                    result[x, y] = -g[x - 1, y - 1] + 0.5 * u[x, y] + 0.5 * u[x - 1, y]
                elif y == 0: # Left edge
                    result[x, y] = 0
                elif y == u.shape[1] - 1: # Right edge
                    result[x, y] = -g[x, y - 1] - g[x - 1, y - 1] + u[x, y] + 0.5 * u[x + 1, y]
                # Inside Image
                else:
                    result[x, y] = u[x, y] + 0.5 * u[x + 1, y] + 0.5 * u[x - 1, y] - g[x, y]
        return result

def derivative_gp_term(u):
    """
    Derivative of Gaussian Prior Regularization Term
    u: The image
    """
    result = np.zeros(u.shape)
    for x in range(u.shape[0]):
        for y in range(u.shape[1]):
            # Corners
            if x == 0 and y == 0:
                result[x, y] = 4 * u[0, 0] - 2 * u[1, 0] - 2 * u[0, 1]
            elif x == 0 and y == u.shape[1] - 1:
                result[x, y] = 4 * u[0, y] - 2 * u[1, y] - 2 * u[0, y - 1]
            elif x == u.shape[0] - 1 and y == 0:
                result[x, y] = 4 * u[x, 0] - 2 * u[x - 1, 0] - 2 * u[x, 1]
            elif x == u.shape[0] - 1 and y == u.shape[1] - 1:
                result[x, y] = 4 * u[x, y] - 2 * u[x - 1, y] - 2 * u[x, y - 1]
            # Edges
            elif x == 0: # Upper edge
                result[x, y] = 6 * u[0, y] - 2 * u[1, y] - 2 * u[0, y + 1] - 2 * u[0, y - 1]
            elif x == u.shape[0] - 1: # Lower edge
                result[x, y] = 6 * u[x, y] - 2 * u[x - 1, y] - 2 * u[x, y + 1] - 2 * u[x, y - 1]
            elif y == 0: # Left edge
                result[x, y] = 6 * u[x, 0] - 2 * u[x - 1, 0] - 2 * u[x + 1, 0] - 2 * u[x, 1]
            elif y == u.shape[1] - 1: # Right edge
                result[x, y] = 6 * u[x, y] - 2 * u[x - 1, y] - 2 * u[x + 1, y] - 2 * u[x, y - 1]
            # Inside Image
            else:
                result[x, y] = 8 * u[x, y] - 2 * u[x - 1, y] - 2 * u[x + 1, y] - 2 * u[x, y - 1] - 2 * u[x, y + 1]

```

```
return result
```

```
def derivative_atv_term(u):
```

```
    """
```

```
    Derivative of Anisotropic Total Variation Regularization Term
```

```
    u: The image
```

```
    """
```

```
    result = np.zeros(u.shape)
```

```
    for x in range(u.shape[0]):
```

```
        for y in range(u.shape[1]):
```

```
            # Corners
```

```
            if x == 0 and y == 0:
```

```
                result[x, y] = - np.sign(u[1, 0] - u[0, 0]) - np.sign(u[0, 1] - u[0, 0])
```

```
            elif x == 0 and y == u.shape[1] - 1:
```

```
                result[x, y] = np.sign(u[0, y] - u[0, y - 1]) - np.sign(u[1, y] - u[0, y])
```

```
            elif x == u.shape[0] - 1 and y == 0:
```

```
                result[x, y] = np.sign(u[x, 0] - u[x - 1, 0]) - np.sign(u[x, 1] - u[x, 0])
```

```
            elif x == u.shape[0] - 1 and y == u.shape[1] - 1:
```

```
                result[x, y] = np.sign(u[x, y] - u[x, y - 1]) + np.sign(u[x, y] - u[x - 1, y])
```

```
            # Edges
```

```
            elif x == 0: # Upper edge
```

```
                result[x, y] = - np.sign(u[1, y] - u[0, y]) - np.sign(u[0, y + 1] - u[0, y])
```

```
                u[0, y] - u[0, y - 1])
```

```
            elif x == u.shape[0] - 1: # Lower edge
```

```
                result[x, y] = np.sign(u[x, y] - u[x - 1, y]) + np.sign(u[x, y] - u[x, y + 1])
```

```
                u[x, y + 1] - u[x, y])
```

```
            elif y == 0: # Left edge
```

```
                result[x, y] = - np.sign(u[x + 1, 0] - u[x, 0]) + np.sign(u[x, 0] - u[x, 1])
```

```
                u[x, 1] - u[x, 0])
```

```
            elif y == u.shape[1] - 1: # Right edge
```

```
                result[x, y] = np.sign(u[x, y] - u[x, y - 1]) + np.sign(u[x, y] - u[x, y + 1])
```

```
                u[x + 1, y] - u[x, y])
```

```
            # Inside Image
```

```
            else:
```

```
                result[x, y] = np.sign(u[x, y] - u[x - 1, y]) - np.sign(u[x + 1, y] - u[x, y])
```

```
                u[x, y + 1] - u[x, y]) - np.sign(u[x, y] - u[x, y - 1])
```

```
    return result
```

```
def GD(g, reg_lambda, action=""): 
```

```
    """
```

```
    Gradient Descent algorithm
```

```
    g: grayscale blurry image of size (M, N)
```

```
    reg_lambda: regularization parameter
```

```
    Optional 'action': Variant to be used.
```

```
    """
```

```
    u = np.zeros((g.shape[0] + 1, g.shape[1] + 1))
```

```
    epsilon = 0.1
```

```
    for x in range(g.shape[0]):
```

```
        for y in range(g.shape[1]):
```

```
            u[x, y] = g[x, y]
```

```
    if action == "gauss": # Gaussian Regularization
```

```
        gradient = derivative_img_term(g, u) + reg_lambda * derivative_gp_term(u)
```

```
        # while (np.sum(np.power(gradient, 2)) > margin):
```

```
        for i in range(1000):
```

```
            u = u - epsilon * gradient
```

```
            gradient = derivative_img_term(g, u) + reg_lambda * derivative_gp_term(u)
```

```
    elif action == "aniso": # Anisotropic Total Variation
```

```
        gradient = derivative_img_term(g, u) + reg_lambda * derivative_atv_term(u)
```

```
        for i in range(1000):
```

```

        u = u - epsilon * gradient
        gradient = derivative_img_term(g, u) + reg_lambda * derivative_atv_term(u)
    else: # No Regularization
        gradient = derivative_img_term(g, u)
        while np.sum(np.power(gradient, 2)) > 0.000001:
            u = u - epsilon * gradient
            gradient = derivative_img_term(g, u)
    return u

```

Linearization and Gauss-Seidel

```

In [6]: def LGS(g, reg_lambda):
        """
        Creating the initial image u (1 unit larger than g in both dimensions), which is set
        Until the difference between the two different u's is smaller than 0.00001, the Line

        g: grayscale blurry image of size (M, N)
        reg_lambda: regularization parameter
        """
        margin = 0.000001
        u = np.zeros((g.shape[0] + 1, g.shape[1] + 1))
        old_u = np.zeros((g.shape[0] + 1, g.shape[1] + 1))

        for x in range(g.shape[0]):
            for y in range(g.shape[1]):
                u[x, y] = g[x, y]

        while np.sum(np.power(old_u - u, 2)) > margin:
            old_u = u
            u = iterative_linearization(g, u, reg_lambda)

        return u

```

```
In [7]: def iterative_linearization(g, u, reg_lambda):
        """
        Gauss Seidel approach with linearized term using sparse matrices (speed-up!).
        g: Grayscale blurry image of size (m, n)
        u: Image passed from LSOR()
        reg_lambda: Regularization parameter
        """
        m, n = u.shape[0], u.shape[1]

        # Flatten image
        vector_u = u.reshape(m * n, 1)

        # Hessian Matrix, following theory from Lec03-04_energy minimization.pdf
        hess_mat = hessian_matrix(u, reg_lambda, k)
        # Compute b
        b = hess_mat * vector_u - (derivative_img_term(g, u) + reg_lambda * derivative_gp_term(u))
        # Triangular lower sparse matrix
        lower_mat = scipy.sparse.tril(hess_mat, k=0)
        # Triangular upper sparse matrix
        upper_mat = scipy.sparse.triu(hess_mat, k=1)

        vector_u = inv(lower_mat) * (b - upper_mat * vector_u)

        # Reshape flat image to initial size
        u = vector_u.reshape(m, n)

        return u
```

```
In [8]: def iterative_linearization_SOR(g, u, reg_lambda, w):
        """
        Successive Over-Relaxation Approach, using sparse matrices (speed-up!) for the linear system.
        g: Grayscale blurry image of size (m, n)
        u: Image passed from LSOR()
        reg_lambda: Regularization parameter
        w: SOR parameter
        """
        m, n = u.shape[0], u.shape[1]

        # Flatten image
        vector_u = u.reshape(m * n, 1)
        # Hessian Matrix
        hess_mat = hessian_matrix(u, reg_lambda, k)
        # Triangular lower sparse matrix
        lower_mat = scipy.sparse.tril(hess_mat, k=-1)
        # Triangular upper sparse matrix
        upper_mat = scipy.sparse.triu(hess_mat, k=1)
        # Diagonal sparse matrix, defined by the three matrices above
        diag_mat = hess_mat - lower_mat - upper_mat

        b = hess_mat * vector_u - (derivative_img_term(g, u) + reg_lambda * derivative_gp_term(u))

        vector_u = inv(diag_mat + w * lower_mat) * (w * b - (w * upper_mat + (w - 1) * diag_mat * vector_u))

        # Reshape flat image to initial size
        u = vector_u.reshape(m, n)

        return u
```

```
In [9]: def LSOR(g, reg_lambda, w):
        """
        Creating the initial image u (1 unit larger than g in both dimensions), which is set
        Until the difference between the two different u's is smaller than 0.00001, the Line

        g: Grayscale blurry image of size (m, n)
        reg_lambda: Regularization parameter
        w: SOR parameter
        """
        u = np.zeros((g.shape[0] + 1, g.shape[1] + 1))
        prior_u = np.zeros((g.shape[0] + 1, g.shape[1] + 1))

        for x in range(g.shape[0]):
            for y in range(g.shape[1]):
                u[x, y] = g[x, y]

        while np.sum(np.power(prior_u - u, 2)) > 0.00001:
            prior_u = u
            u = iterative_linearization_SOR(g, u, reg_lambda, w)
        return u
```

Parameters

```
In [10]: reg_lambda = 0.005 # lambda regularization parameter. you need to play with this
        # 0.025 best for Gauss
        # 0.005 best for anisotropic total variation

        # k matrix. Look below how to choose it
        k = 1
```

$$k = 0 \longrightarrow \begin{bmatrix} \frac{1}{2} & \frac{1}{2} \\ 0 & 0 \end{bmatrix} \quad k = 1 \longrightarrow \begin{bmatrix} \frac{1}{2} & 0 \\ \frac{1}{2} & 0 \end{bmatrix} \quad k = 2 \longrightarrow \begin{bmatrix} \frac{1}{2} & 0 \\ 0 & \frac{1}{2} \end{bmatrix} \quad k = 3 \longrightarrow \begin{bmatrix} 0 & \frac{1}{2} \\ \frac{1}{2} & 0 \end{bmatrix}$$

Blurring the image

```
In [11]: def blurring(img, kernel_mat):
        """
        Blur image with convolution, could alternatively be done simply with scipy.signal.convolve2d
        img: The image passed
        kernel_mat: The kernel used
        """
        blurred_img = np.zeros((img.shape[0] - 1, img.shape[1] - 1))
        for row in range(img.shape[0] - 1):
            for column in range(img.shape[1] - 1):
                blur_pixel = 0
                for p in range(kernel_mat.shape[0]):
                    for q in range(kernel_mat.shape[1]):
                        blur_pixel += kernel_mat[p, q] * img[row - p + 1, column - q + 1]
                blurred_img[row, column] = blur_pixel
        return blurred_img
```

Miscellaneous functions

```
In [12]: def compute_loss(blurred_img, deblurred_img, kernel_mat):
    """
    Computes the data loss
    blurred_img: The blurred image
    deblurred_img: The deblurred image
    kernel_mat: The kernel used
    """
    loss = 0
    for x in range(blurred_img.shape[0]):
        for y in range(blurred_img.shape[1]):
            convolution = 0
            for p in range(kernel_mat.shape[0]):
                for q in range(kernel_mat.shape[1]):
                    convolution += kernel_mat[p, q] * deblurred_img[x - p + 1, y - q + 1]
            loss += (blurred_img[x, y] - convolution) ** 2
    return loss

def compute_diff(deblurred_img, sample_img):
    """
    Calculates the sum of squared differences
    deblurred_img: The deblurred image
    sample_img: The unblurred image
    """
    return np.sum(np.power(deblurred_img - sample_img, 2))
```

Reading image

You are provided with only the synthetic image. The purpose of using a synthetic image is to see possible changes in the deblurred image after applying your algorithm. You can convolve the original image using the [scipy library \(https://docs.scipy.org/doc/scipy/reference/generated/scipy.signal.convolve2d.html\)](https://docs.scipy.org/doc/scipy/reference/generated/scipy.signal.convolve2d.html). Note that you must choose **valid** as the mode in your convolution.

```
In [13]: img_org = sample_image = create_random_binary_img_1()

# Kernel k = 1
kernel = np.array([[0.5, 0], [0.5, 0]])

img_blurred = blurring(sample_image, kernel)
```

GRADIENT DESCENT

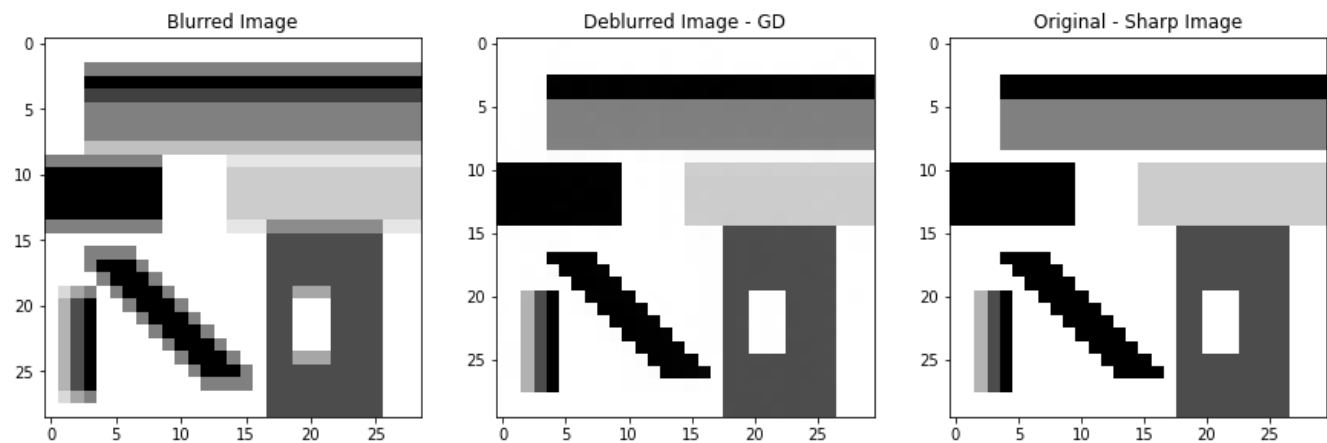

```
In [14]: # run gradient descent algorithm
img_deblurred_gd = GD(img_blurred, reg_lambda, "aniso")

# for visualization purpose
# img_deblurred_gd = np.clip(u_final_gd, 0, 1)

plt.figure(figsize=(15, 5))
ax1 = plt.subplot(1, 3, 1)
plt.imshow(img_blurred, cmap='gray')
ax2 = plt.subplot(1, 3, 2)
plt.imshow(img_deblurred_gd, cmap='gray')
ax3 = plt.subplot(1, 3, 3)
plt.imshow(img_org, cmap='gray')

ax1.set_title("Blurred Image")
ax2.set_title("Deblurred Image - GD")
ax3.set_title("Original - Sharp Image")
plt.show()

print(f"For lambda {reg_lambda:.5f}\t \nk_type {k:1d}")
print(f"Loss Data Term = {compute_loss(img_blurred, img_deblurred_gd, kernel):.5f}")
print(f"Sum of Squared Distance: deblurred vs. original = {compute_diff(img_deblurred_gd, img_org):.5f}")
```



For lambda 0.00500

k_type 1

Loss Data Term = 0.00733

Sum of Squared Distance: deblurred vs. original = 0.01874

SOR

```

In [15]: w = 1.075 #play with this

# run LSOR algorithm
img_deblurred_sor = LSOR(img_blurred, reg_lambda, w)

plt.figure(figsize=(15, 5))
ax1 = plt.subplot(1, 3, 1)
plt.imshow(img_blurred, cmap='gray')
ax2 = plt.subplot(1, 3, 2)
plt.imshow(img_deblurred_sor, cmap='gray')
ax3 = plt.subplot(1, 3, 3)
plt.imshow(img_org, cmap='gray')

ax1.set_title("Blurred Image")
ax2.set_title("Deblurred Image - LSOR")
ax3.set_title("Original - Sharp Image")
# plt.savefig("results_" + str(k) + ".png")
plt.show()

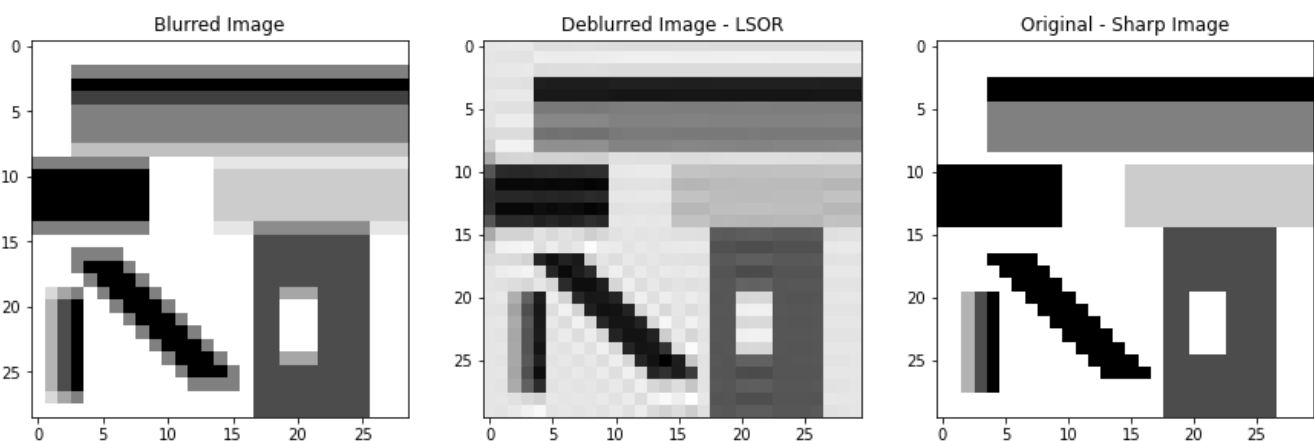
print(f"For lambda {reg_lambda:.5f}")
print(f"Loss Data Term = {compute_loss(img_blurred, img_deblurred_sor, kernel):.5f}")
print(f"Sum of Squared Distance: deblurred vs. original = {compute_diff(img_deblurred_sor, img_org):.5f}")

```

```

c:\users\brian\appdata\local\programs\python\python39\lib\site-packages\scipy\sparse\linalg\dsolve\linsolve.py:318: SparseEfficiencyWarning: splu requires CSC matrix format
warn('splu requires CSC matrix format', SparseEfficiencyWarning)
c:\users\brian\appdata\local\programs\python\python39\lib\site-packages\scipy\sparse\linalg\dsolve\linsolve.py:215: SparseEfficiencyWarning: spsolve is more efficient when sparse b is in the CSC matrix format
warn('spsolve is more efficient when sparse b '

```



```

For lambda 0.00500
Loss Data Term = 0.03840
Sum of Squared Distance: deblurred vs. original = 1.93167

```

GAUSS - SEIDEL

```

In [16]: nb_iter = 1

# run LGS algorithm
img_deblurred_gs = LGS(img_blurred, reg_lambda)

plt.figure(figsize=(15, 5))
ax1 = plt.subplot(1, 3, 1)
plt.imshow(img_blurred, cmap='gray')
ax2 = plt.subplot(1, 3, 2)
plt.imshow(img_deblurred_gs, cmap='gray')
ax3 = plt.subplot(1, 3, 3)
plt.imshow(img_org, cmap='gray')


ax1.set_title("Blurred Image")
ax2.set_title("Deblurred Image - LGS")
ax3.set_title("Original - Sharp Image")
#plt.savefig("figures/results_LSOR_"+str(k)+".png")
plt.show()

print(f"For lambda {reg_lambda:.5f}")
print(f"Loss Data Term = {compute_loss(img_blurred, img_deblurred_gs, kernel):.5f}")
print(f"Sum of Squared Distance: deblurred vs. original = {compute_diff(img_deblurred_gs, img_org):.5f}")

```

c:\users\brian\appdata\local\programs\python\python39\lib\site-packages\scipy\sparse\linalg\dsolve\linsolve.py:144: SparseEfficiencyWarning: spsolve requires A be CSC or CSR matrix format

warn('spsolve requires A be CSC or CSR matrix format',



For lambda 0.00500

Loss Data Term = 0.03844

Sum of Squared Distance: deblurred vs. original = 1.92149

Sum of Squared Distances Calculation

```
In [17]: lambdas = [0, 0.0001, 0.0005, 0.001, 0.005, 0.01, 0.05, 0.1, 0.5, 1]
ssd_atv = []
ssd_gauss = []

for value in lambdas:
    img_anisotropic = GD(img_blurred, value, "aniso")
    ssd_atv.append(compute_diff(img_anisotropic, sample_image))

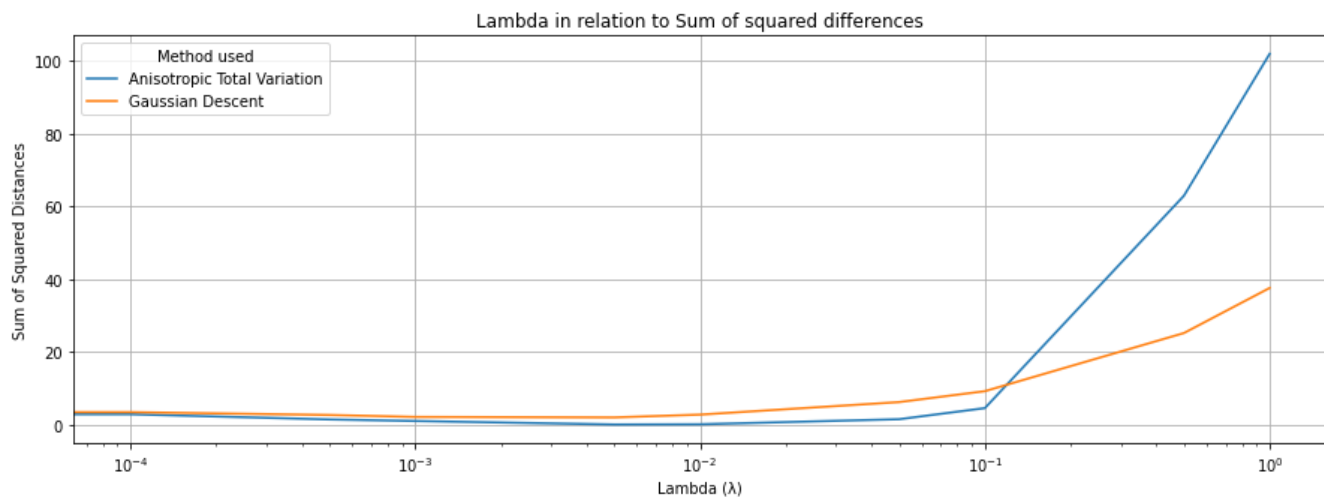
    img_gaussian = GD(img_blurred, value, "gauss")
    ssd_gauss.append(compute_diff(img_gaussian, sample_image))
```

Sum of Squared Distances Plotting

```
In [19]: plt.figure(figsize=(15, 5))

plt.title("Lambda in relation to Sum of squared differences")

plt.xlabel("Lambda ( $\lambda$ )")
plt.ylabel("Sum of Squared Distances")
plt.plot(lambdas, ssd_atv, label = "Anisotropic Total Variation")
plt.plot(lambdas, ssd_gauss, label = "Gaussian Descent")
plt.legend(title='Method used')
plt.xscale("log")
plt.grid()
plt.show()
```



Optimal values are between 10^{-3} and 10^{-1} as expected, thanks to the exercise description. This fits for both Anisotropic Total Variation and Gaussian Descent.

If we move above the optimal values, the reconstructed image is very noise and below the impact of the regularization term is negligible.